

## Assertion

- General constraints on DB.

- Create Assertion name  
check (<condition>);

week 3

ex:

C. A. myAssert check

(NOT EXISTS (

SELECT p.name

FROM p, pr

where p.name = pr.name

group by p.name

having Count(\*) > 200))

→ no item should be sold  
more than 200 times

Whenever a modification happens (deletion, insertion, update), assertions  
will be checked. Not practical.  
on the table/relation

ex:

C. A. Fewer check (

(Select Count(\*) from bars) <=

(Select Count(\*) from drinkers))

→ There cannot be more bars  
than drinkers.

For better performance, deletion on bars and insertion on drinkers should  
not invoke this assertion. Therefore, use triggers.

## Trigger

event → ins, up, del. to the relation

condition → can be expressed as either before or after the event

if condition true

action

## row trigger

- exe only for only modified tuple

ex: When price is ↓, set category 'on sale'

C. T. Productcategories

after update of price on Product  
referencing

new row as newtuple  
old row as oldtuple

event

for each row

when ( Old.price > new.price ) ] condition

Update Product

set cat = 'on sale'

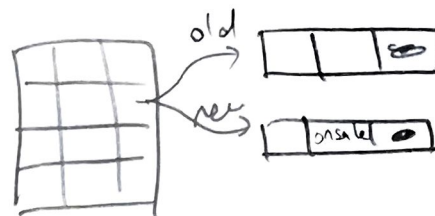
where pid = oldtuple.pid

action

- DB will hold the table in the ram.

- When price is changed it knows

- It keeps copy of old and new row before  
storing it permanent



## Statement trigger

after update

referencing

new table as —

old table —

begin

update product

end

Instead of F.K., use triggers to insert a beer without manu.  
Sells (beer, beer, price) → if there is FK beer, then I cannot insert  
Beer (beer, manu) beer if there is no entry in Beer table

C. T. Beertrip  
after insert on Sells  
ref. new row as newtuple  
for each row  
when (new.beer not in (  
select name from beers))  
insert into beers (name)  
values (new.beer);

insert into sells values (---) → 1 insert  
insert into sells select \* from --- → Bulk insertion (more than 1)  
→ that's why we use → referencing new row as newrow  
# for each row \*

- Drop Product → delete table from db
- delete Product → delete rows in product table
- drop Purchaseprice → we can drop a trigger
- alter Product add column ---
- create or replace trigger --- → if trigger exists, update  
else create

whether you have after or before in trigger, the modification  
will not be stored until trigger finishes.

after update of price on sells → if we say after update on sells, it is not very efficient

- You cannot write trigger inside another trigger. But you can design a DB where trigger causes a sequence of rules. Dangerous.

after delete —  
old row — ✓  
~~new row~~

after insert —  
~~old row~~  
new row ✓

- Empty tables will evaluate true. Empty table does not violate assertion or trigger.

## Views

- For representing different info to diff users.

- It physically does not exist ★

C. Views CP as → virtual table

→ Virtual view

- Computed on demand, slow at runtime

- always up to date

Materialized View

- Pre computed, fast on runtime

- May have stale data

- In big data warehouses, it (material view) is stored in the harddrive (physically)

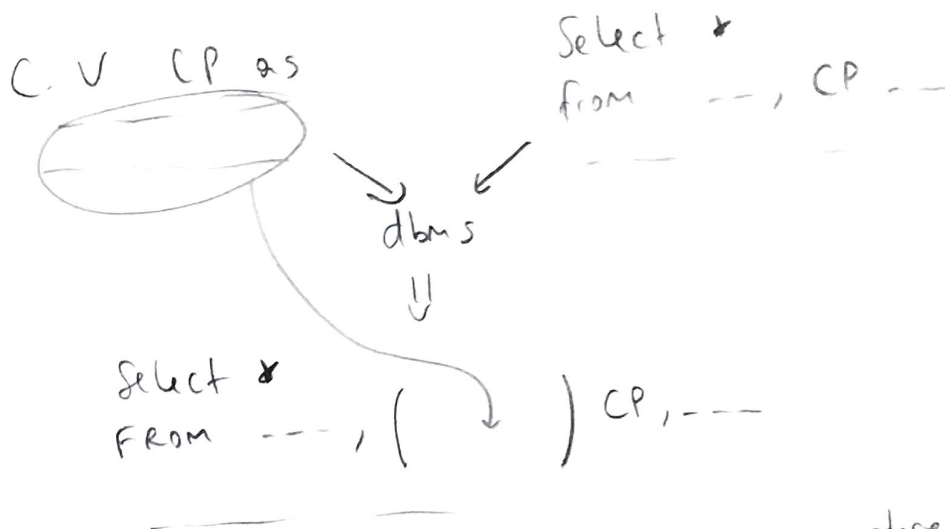
Create [Materialized] View

↓

optional

default is virtual.

- DBMS will unnest the view into query



ex: C. V. Cust.Pr. as

S x.customer, y.price  
F Pur x, Pro y  
W x.pro = y.pname

S u.customer, v.store  
F CustPr u, Purc v  
W u.cust = v.cust and  
u.price > 100

⇒ ①

S u.cust, v.store  
F ( S x.cust, y.price  
F Pur x, Pro y  
W x.pro = y.pname ) u, Purc v  
W u.cust = v.cust and  
u.price > 100

⇓ ② modification

S x.cust, v.store  
F Pur x, Pro y, Pur v  
W x.cust = v.cust and  
y.price > 100 and  
x.product = y.pname

## Applications

- Physical data independence  
↳ vertical / horizontal part
- Security (who can see which table)

# Vertical Partitioning

Resumes

SSN	name	add.	res	pic

T1

SSN	name	add.

T2

SSN	resume

T3

SSN	picture

C.V Resumes AS

Sel T1.SSN, T1.name  
 From T1, T2, T3  
 W T1.SSN = T2.SSN and T2 ---

→ Assume we designed db s.t.  
 we have T1, T2 and T3.  
 Now, we want to combine them  
 for a view

The query



S.	add
F	Resumes
W	name = 'Sue'

S add  
 F ( S SSN, name, add. ---  
 F T1, T2, T3  
 W T1.SSN = T2.SSN ---  
 W name = 'Sue'

If the dbms is clever, it will optimize the query above. But, not all dbms' are clever

When to do vertical part?

- Some fields are large and rarely accessed → picture
- distributed db
- T1 comes from 1 source, T2 comes from another



# Horizontal Partitioning

customers

SSN	name	city	Country
		Houston	USA
		Seattle	USA
		Portland	Canada

C in Houston

SSN	name	city	Country
		Houston	

C in Seattle

1	2	3	4
		Seattle	

C in Canada

1	2	3	4
			Canada

Assume that we have C in Houston, C in Seattle ---  
and we want to combine them.

C. V. Customers as  
C in Houston  
UNION all  
C in Seattle

→ if we say

SELECT name  
FROM Customers  
WHERE city = 'Seattle'

DBMS cannot optimize this.  
It will bring all data  
then eliminate non  
Seattle. Too much  
work.

However, if we create our view in this format

C V Customers

(SELECT \*  
FROM Customers  
WHERE city = 'Seattle')

UNION all

(SELECT \*  
FROM Customers  
WHERE city = 'Houston')

Compiler can optimize  
the query above  
Before doing execution.

## Views And Security

C

name	addr	balance

Some users are not allowed to see this column

C. V Public C.

S name, addr

F Customers



C

name	addr	balance
		+450
		+20
		-30

Some are not allowed to see >0 balance

C. V --

S \*

F Cust W Balance < 0

## Privileges

Select, Insert, Delete, Update

row based

may apply to only one attribute

Grant <List of privileges>  
ON <relation or objects>  
TO <db user>

Grant Sel, Upd (address)  
ON PubCust  
TO Joe

Joe can select and update the address of PubCust

## Updating Views

- A view update changes underlying table → produce requested change to the view.



Transcript (Studid, Ciscode, Semester, Grade)

C.V. CorgReg (Studid, ciscode, semester) as

S T.studid, T.ciscode, T.semester

F Tran. T

W T.ciscode Like 'Corg%' and T.Semester = 'S2013'

Q: insert into CorgReg (---) vals (1111, Corg352, S2013)  
What will be the grade?

A: In transcript table grade will be NULL (if it is allowed)

Q: insert into CorgReg (---) vals (1111, Econ210, S2013)

New tuple will not be in the view.

A: Allow insertion WITH CHECK OPTION clause on the Create View

Professor (Id, name, dept)

Dept (deptid, name)

C.V. ProfDept (P.name, D.name) as

S P.name, D.name

F P.P, D.D

W P.deptid = D.id

Q: delete <Smith, Corg> From ProfDept

1) Should professor be deleted??

2) Should Corg dept be deleted??

3) Update row for smith deptid=null??

↳ Good, but computer cannot know this.

Ambiguity

- We cannot write Foreign Keys in Views for solving ambiguity.
- But we can write Triggers!

## Some Views are not Updatable

C.V AvgSalary (deptId, avg sal) as  
 S E.deptId, AVG (E.salary)  
 F Emp E  
 Group By E.deptId

How do we handle

(update AvgSalary  
 Set AvgS = 1.1 \* Avg sal) ?

Employee (ssn, name, dept, project, salary)

C.V. Developers as  
 S name, project  
 F Emp  
 W dept = 'Development'

→ insert into developers vals (Joe, 'optimizer')

⇓  
 insert into Employee vals  
 (null, Joe, null, optimizer, null)

↓  
 assume it  
 is allowed

↓  
 this must  
 be development

To solve the issue above, we can use triggers.

C.T AllowInsert

**Instead OF** insert on Developers

Referencing new row as newrow

For each row

Begin

insert into emp (name, dept, proj)

vals (newrow.name, 'Development', newrow.project)

End

## Materialized Views

- Each time base table change, mat view may change.
- We cannot afford to recompute the view with each change.
- Soln: Periodically construct mat view.

## Indexes

- Hash table or B+tree (default)

Syntax → Create Index C on Product (cat)  
C I Emplnd on Emp (name) USING HASH;  
— — — Trnsript (semester, course) USING BTREE;  
                    ↓                    ↓  
                  1<sup>st</sup> semester      2<sup>nd</sup> course

Q: Find prices of beers manuf by Tuborg and sold at Joe's bar

C. I. BeerInd on Beers (manuf)  
C - I. SellInd on Sells (beer, beer)

Query  
S price  
F Beers, Sells  
WH manuf = Tuborg and  
B.name = S.beer and  
bar = Joe's bar

→ much faster.

★ If insert, delete, update on the index is very frequent, these operation will be slow.