

# TRANSACTIONS

week 6

- sequence of actions
- either all or none.

Autocommit off  $\rightarrow$  first SQL query starts txn

Start transaction

[SQL]

COMMIT or ROLLBACK (=ABORT)

if app cannot continue because of integrity cond., dbms will rollback

ABORT reasons

- $\rightarrow$  ctrl+c
- $\rightarrow$  Condition fails
- $\rightarrow$  deadlock (system can abort)

Without transactions

## Lost Updates

Upd Cust

Set rental += 1

Where name = "Fred"

Q1

Same query

in the left

Q2

We want at the end  
rental  $\rightarrow$  rental + 2 but  
it might be +1.

## Unrepeatable Read

x = Sel. rentals

if (x < 5) {

    upd Cust

    set rentals += 1

}

in here, it might read 4.

However, another txn might set it to 5

here

Then, in this statement, we will have rentals = 6 instead of 5.

This read should be 4 again, but it is updated.

## Inconsistent Read

Upd Pro.  
Set Quantity += 5  
Where pro = guitar

select sum(quantity)  
from Product.

Here, we will read 5 more.

upd Pro  
set Quantity -= 5  
Where product = guitar

## Dirty Read

Client 1 100\$ acc1 → acc2

X = acc1.balance

acc2.balance += 100

if (X > 100) acc1.balance -= 100

else // Rollback  
acc2.balance -= 100

Client 2 100\$ acc2 → acc3

Y = acc2.balance

acc3.balance += 100

if (Y > 100) acc2.balance -= 100

else // Rollback  
acc3.balance -= 100

Dirty read  $\cong$  Inconsistent read

## Crash

upd acc  
set bal -= 100  
none free

CRASH

upd acc  
set bal += 100  
none freed

# ACID

Atomicity → all or none

Consistency → our responsibility  
 ↳ writing correct txn  
 ↳ integrity rules

Isolation → concurrency

Durability → will survive after failure (write to disk)  
 ↳ by dbms

## Recap:

Dirty Reads

<u>T<sub>1</sub></u>	<u>T<sub>2</sub></u>
W(A)	R(A)

Write-read conflict

Abort

Inconsistent Read

<u>T<sub>1</sub></u>	<u>T<sub>2</sub></u>
A=20, B=20 W(A)	R(A) R(B)

→ it will not read B  
write-read conflict

W(B)

Unrepeatable Read

<u>T<sub>1</sub></u>	<u>T<sub>2</sub></u>
W(A)	R(A)
	R(A)

read-write conflict

Lost Update

<u>T<sub>1</sub></u>	<u>T<sub>2</sub></u>
R(A)	R(A)
A+5 W(A)	A+1.3 W(A)

Write-write conflict

→ lost

# Isolation

## Serial

T1  
A=100  
B=100

T2

A=2  
B=2

## Serializable

A schedule is serializable, if it is equivalent to a serial schedule  
This is serializable schedule  
not serial schedule

T1  
R(A, t)  
t+=100  
W(A, t)

T2

R(A, s)  
S=2  
W(A, s)

R(B, t)  
t+=100  
W(B, t)

R(B, s)  
S=2  
W(B, s)

## Non Serializable

T1  
R(A)  
A+=100  
W(A)

T2

R(A)  
~~2~~  
W(A)  
R(B)  
B=2  
W(B)

A=100, B=10

A=100

A=200

B=20

B=120

They both should be 200.

R(B)  
B+=100  
W(B)

## edge case

T1

R(A)  
A=100  
W(A)

T2

R(B)  
B=200  
W(B)  
R(A)  
A=200  
W(A)

A=10 B=10

A=310 B=310

it is serializable schedule  
- T1 then T2 is equivalent to this schedule.  
- However, we don't expect scheduler to find this.

R(B)  
B=100  
W(B)

## Conflict Serializability

### Conflicts

$r_i(x), w_j(y)$  → Two actions by same txn.

$w_i(x), w_j(x)$  → Two writes on same element by diff txn

$w_i(x), r_j(x)$   
 $r_i(x), w_j(x)$  → read/write on same element by diff txn

Defn: A schedule is conflict serializable, if it can be transformed into a serial schedule by a series of swaps of adjacent non-conflicting actions

ex:  
 $r_1(A) w_1(A) r_2(A) w_2(A), r_1(B) w_1(B) r_2(B) w_2(B)$

$r_2(A) r_1(B), w_2(A)$

$r_1(B) r_2(A), w_2(A) w_1(B)$


⇓

$r_1(A) w_1(A) r_1(B) w_1(B), r_2(A) w_2(A) r_2(B) w_2(B)$

# Use precedence graph


ex:

$r_2(A) r_1(B) w_2(A) r_3(A) w_1(B) w_3(A) r_2(B), w_2(B)$




ex:

$r_2(A) r_1(B) w_2(A) r_2(B) r_3(A) w_1(B) w_3(A) w_2(B)$




## View Equivalence

Schedules



$w_1(X) w_2(X) w_2(Y) w_1(Y) w_3(Y)$



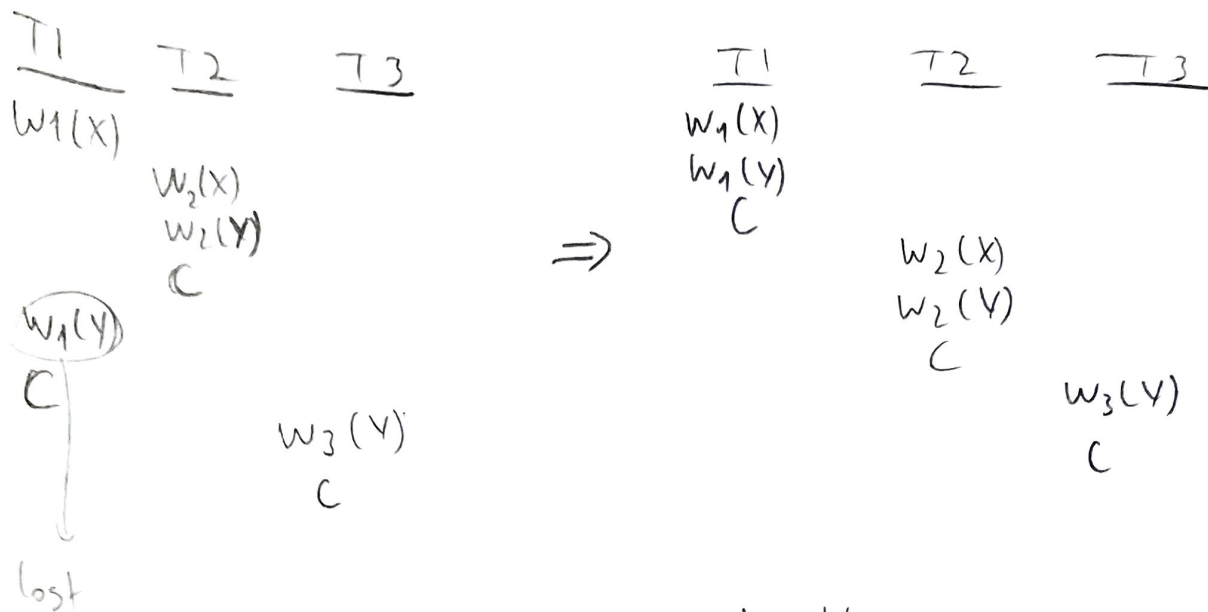
this write will be lost, because  $T_1$  and  $T_3$  will overwrite.

equivalent  $\rightarrow w_1(X) w_1(Y) w_2(X) w_2(Y) w_3(Y)$

not conflict equivalent



ex:




Serializable, but not conflict serializable.


Two schedule is view equivalent:

S  
reads A initially

S'  
reads A initially  $\rightarrow$  initial reads



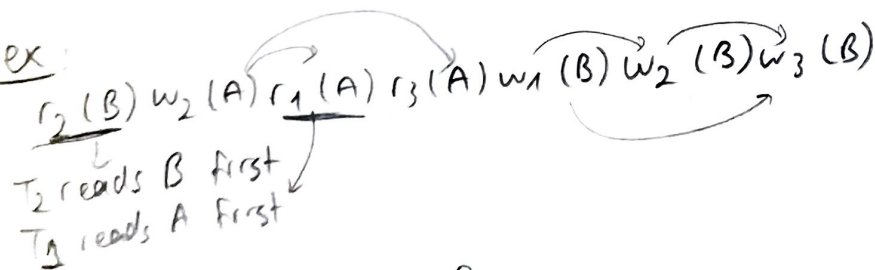
=



$\rightarrow T_1$  reads  $T_2$ 's write

$\rightarrow$  final writes

ex:




A

$T_2$  writes  
 $T_1$  reads  
 $T_3$  reads

B

$T_2$  reads

$\rightarrow$  A'ya en son  $T_2$  yazıyor.



$T_1$  writes  
 $T_2$  writes  
 $T_3$  writes

$T_1 - T_3$  sırası nasıl olmalı?


$\rightarrow$  en son  $T_3$  yazdığı için  $T_3$  olmalı

$T_2 - T_1 - T_3$

$r_2(B) w_2(A) w_2(B) r_1(A) w_1(B) r_3(A) w_3(B)$




view equivalent  $\rightarrow$

## Question



## Recoverable Schedules

- We have to read transaction that have committed.
- Conflict serializable



## Review

### Serializability

- Serial
- Serializable
- Conflict serializable
- View serializable

### Recoverability

- Recoverable
- Avoid cascading aborts



# Locks

-Pessimistic scheduler

$L_i(A) \rightarrow TX_i$  acquires lock for A

$U_i(A) \rightarrow TX_i$  releases lock for A

## Non serializable schedule

T1  
R(A)  
A+100  
w(A)

T2

R(A)  
A\*2  
w(A)  
R(B)  
B\*2  
w(B)

$\Rightarrow$

T1

$L_1(A)$   
R(A)  
A+100  
w(A),  $U_1(A)$   $L_1(B)$

R1(B)  
B+100  
w(B)  $U_1(B)$

T2

$L_2(A)$   $R_2(A)$   
A\*2  
w2(A)  $U_2(A)$   
 $L_2(B)$  // DENIED, wait

// Granted

R(B)  
B\*2  
w(B)  $U_2(B)$

conflict serializable ✓

wrong case

T1

$L_1(A)$   
R(A)  
A+100  
w(A)  
 $U_1(A)$

T2

$L_2(A)$   $R_2(A)$   
A\*2  
w(A)  $U_2(A)$   
 $L_2(B)$   $R_2(B)$   
B\*2  
w(B)  $U_2(B)$

not conflict serializable!  
Locks did not enforce that!

$L_1(B)$   $R_1(B)$   
B+100  
w(B)  $U_1(B)$

# Two Phase Locking (2PL)

- In every txn, all lock request must precede all unlock



ex:


T<sub>1</sub>  
 $L_1(A) L_1(B)$   
 $R(A)$   
 $A++$   
 $w(A) U_1(A)$

T<sub>2</sub>  
 $L_2(A) R(A)$   
 $A * 2$   
 $w(A)$   
 $L_2(B)$  // denied, it waits until the lock is released

$U_2(B)$  // granted

Theorem 2PL ensures conflict serializability

Proof: Suppose it does not ensure, then, there is cycle in the graph.



Then, there is temporal cycle in the schedule

$U_1(A) \rightarrow L_2(A)$  // 1st txn, A's b/w txn 2 about  
 $L_2(A) \rightarrow U_2(B)$   
 $U_2(B) \rightarrow L_3(B)$   
 $L_3(B) \rightarrow U_3(C)$   
 $U_3(C) \rightarrow L_1(A)$

## 2PL and nonrecoverable Problem


T<sub>1</sub>  
 $L_1(A) U_1(B)$   
 $==$   
 $==$   
abort

T<sub>2</sub>  
 $L_2(A)$   
 $L_2(B)$  // denied  
 $==$   
commit

Problem!

## Strict 2PL

- Locks are released at COMMIT or ROLLBACK
- ensures recoverable
- avoids cascading aborts



No shrinking

T<sub>1</sub>  
L<sub>1</sub>(A) R(A)  
W(A)

T<sub>2</sub>  
 L<sub>2</sub>(A) // denied

L<sub>1</sub>(B) R(B)  
W(B)  
W<sub>1</sub>(A) W<sub>1</sub>(B)  
 Rollback ✓ // Granted ✓

- Strict schedule does not allow dirty reads and dirty writes.
- It can still be deadlock!

T<sub>1</sub>  
L<sub>1</sub>(A) R(A)  
W(A)

T<sub>2</sub>  
L<sub>2</sub>(B) R(B)  
L<sub>2</sub>(A) // denied  
L<sub>2</sub>(B) // denied



## Issues

- Implementation
- Lock modes
- Granularity
- Deadlocks
- Performance (looks like serial schedules)

# Locking Scheduler

Task 1 → behalf of txn  
 ↳ read/write, lock request, strict 2PL

Task 2 → behalf of system  
 ↳ Lock table, when a lock is requested check the lock table.  
 ↳ when txn abort, release all locks  
 ↳ check for deadlocks occasionally.

Fine granularity → rows, tuples  
 ↳ high concurrency  
 ↳ #locks ↑, high overhead

Coarse grain → tables, predicate locks  
 ↳ less concurrency  
 ↳ less overhead  
 ↳ many false conflict (we think that it is locked, it is being used) but it is not

SQL  
 From R1  
 where R1.0 > 100

## Lock Modes

S → shared → read → more than 1 txn can read  
 X → exclusive → write → only 1 txn can write, none can read/write

## Hierarchical Locking

- To place a lock on element, start at the top.

- If element to lock, get S or X.

- If element is in deeper level, leave internal lock, IS, IX

To get	Must have all ancestors
IS or S	IS or IX
IX, SIX, X	IX or SIX

SIX → I will read all table, update some of them in some condition.

<u>page</u>	<u>Lock</u>	<u>current</u>	<u>wait</u>
P1	S	{T <sub>1</sub> }	{T <sub>2</sub> }

T<sub>1</sub> was reading P1. T<sub>1</sub> changed its mind and want to write on P1.

lock upgrade

P1	X	{T <sub>1</sub> }	{T <sub>2</sub> }
----	---	-------------------	-------------------

### Problem

S<sub>1</sub>(A) S<sub>2</sub>(A) X<sub>2</sub>(A) X<sub>1</sub>(A)

P1	S	{T <sub>1</sub> , T <sub>2</sub> }	{ }	S <sub>1</sub> (A), S <sub>2</sub> (A)
P1	S	{T <sub>1</sub> , T <sub>2</sub> }	{T <sub>2</sub> }	X <sub>2</sub> (A)
P1	<del>S</del>	{T <sub>1</sub> , T <sub>2</sub> }	{T <sub>1</sub> , T <sub>2</sub> }	X <sub>1</sub> (A)

dead lock!