

Cost = read, not write

$B(R)$ = # of block (pages) for R

$T(R)$ = # of tuples in R

$V(R, a)$ = # of distinct values of attribute a

when a is key
else

$$V(R, a) = T(R)$$

$$V(R, a) < T(R)$$

Join ($R \bowtie S$)

1) One-pass Algorithms

It fits in the memory

a) Hash Join



$R \bowtie S$

- scan R , build buckets in mem.


- scan S / join.

Cost $B(R) + B(S)$.

$M = 2L = (\text{large enough})$



read and hash patient



2-2 match, write to output.

repeat until all S finishes

each block contains 2 tuple (assumption)
(page)

Joined record, output buffer contains:

(pid, name, addr, provider, policy)

only 1 pid (They have been joined on pid)

Cost
$B(R) + B(S)$

Algorithm RMS

Pipelined.

```

Open() {
    H = newHashTable();
    R.open();
    x = R.getNext();
    while (x != null) {
        H.insert(x);
        x = R.getNext();
    }
    R.close();
    S.Open();
    buffer = [];
}

```

```

GetNext() {
    while (buffer == []) {
        x = S.getNext();
        if (x == null) return null;
        buffer = H.find(x);
    }
    z = buffer.first();
    buffer.reset();
    return z;
}

```

Since it is pipelined
1 element at a time.

b) Nested loop join

for each tuple t_1 in R
 for each tuple t_2 in S
 if t_1 and t_2 join then output (t_1, t_2)

Cost
$B(R)$
+
$T(R) \cdot B(S)$

read S
of tuples of R times


c) Page at a time

for each page of $R \rightarrow B(R)$
 for each page of $S \rightarrow B(R) \cdot B(S)$
 for all pairs of tuples t_1 in R, t_2 in S
 if t_1 and t_2 join, output (t_1, t_2)

Cost
$B(R) + B(R) \cdot B(S)$

②

Visualization of page at a time



d) Block Nested loop

Memory has M pages. Use 1 for input
1 for output.

read $(M-2)$ pages of R to memory.

for each group of $(M-2)$ pages r in R : $\rightarrow B(R)$


for each page of tuples s in S : $\rightarrow \frac{B(R)}{(M-2)} \cdot B(S)$

for all pairs of tuples t_1 in r , t_2 in s
output

Cost
$B(R) + \frac{B(R)}{(M-2)} \cdot B(S)$

ex:

$M=6$



if R is 4 pages.
I can fit R into mem ~~4~~

$$\frac{B(R)}{4} + \frac{B(R)-4}{M-2-4} \rightarrow 1 \cdot B(S)$$

if R is 8,

$$\frac{B(R)}{8} + \frac{B(R)-8}{M-2-4} \rightarrow 2 \cdot B(S)$$

I need to traverse B 2 times.

(3)


e) Sort - Merge

Scan R, sort in mem
Scan S, //
Merge R & S.

cost: $B(R) + B(S)$ ← if it fits mem
typically this is not one pass alg.

2) Index Based Algorithms

- Existing index on join attribute (inner)



Selection on equality: $\sigma_{a=v}(R)$ ← Before join

Index based selection

Clustered index on a: $\frac{B(R)}{V(R,a)}$
(ignore I/O for index access)

Unclustered index on a: $\frac{T(R)}{V(R,a)}$

ex: $B(R) = 2000$

$T(R) = 100,000$

$V(R,a) = 20 \rightarrow 20$ distinct vals on a.

$\sigma_{a=v}(R) = ?$

a) Table scan

$B(R) = 2000$ I/O

b) Clustered index $\frac{B(R)}{V(R,a)} = 100$

c) Unclustered index $\frac{T(R)}{V(R,a)} = 5000$



each will have 5000 tuple pointers.

2000 page ver. 20 farklı a ver. ortalaması $\frac{2000}{20} = 100$ page'de bul $|a=v|$ eşitliği vardır. Index clustered olduğundan ilk olan bul, sonra linear search.

indexten hangi page olduğuna bulduk ama unclustered



in the worst case you need 5000 I/O.

④

Dont build unclustered indexes when $V(R, a)$ is small.


end of index based selection

a) Index Nested loop join

$R \bowtie S$ → assume inner (S) has index on join attribute.

Iterate over R, for each tuple, fetch corresponding tuples from S.

Product \bowtie pid Purchase



c1 p1
c2 p1
c1 p2
— p3

→ clustered index on pid

① Find matching in S.

② next in Product.

$B(R) \rightarrow$ read outer relation.

$$B(R) + T(R) \cdot \frac{B(S)}{V(S, a)}$$

↓
for each row
of the R

→ dıyelim 30 page var.
3 distinct a value'ler var.
Bu a value'leri 10 tane page de'ler.
 $T(R) = 10 \rightarrow$ here oku.

$$\text{Clustered index} = B(R) + T(R) \cdot \frac{B(S)}{V(S, a)}$$

$$\text{Unclustered index} = B(R) + T(R) \cdot \frac{T(S)}{V(S, a)}$$

→ unclustered oldıysa
ıfın bütın row lara
bakarak 2 gerekebilir.

3) Two Pass Algorithms

Files are big, dont fit in the memory.

- sorting & hashing.

a) Sort & Merge

External Sorting

- Sort a file $B \rightarrow$ size with mem M .

- Sorting is two pass when $B < M^2$


ex: Page size 32KB

mem 32GB
 $= 10^6$ pages

R can be large as 10^{12} pages


$$10^{12} \times 32 \times 10^3 = 32 PB$$

Step 1:



Step 2

merge $M-1$ runs to a new run



$$? \quad M(M-1) \cong M^2 \quad ?$$

Cost: Read + Write + Read (final write is not included) $\rightarrow 3B(R)$

Assumption $B(R) \leq M^2$

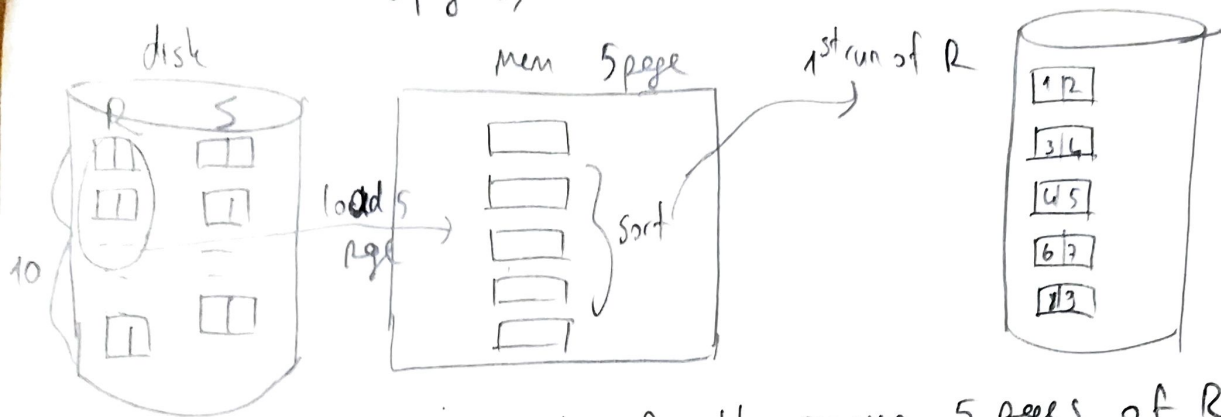
Merge Join

1a: Generate initial runs of 12

1b: // \leq

2: Merge & join \rightarrow either \rightarrow 1st merge then join
 or \rightarrow at the same time

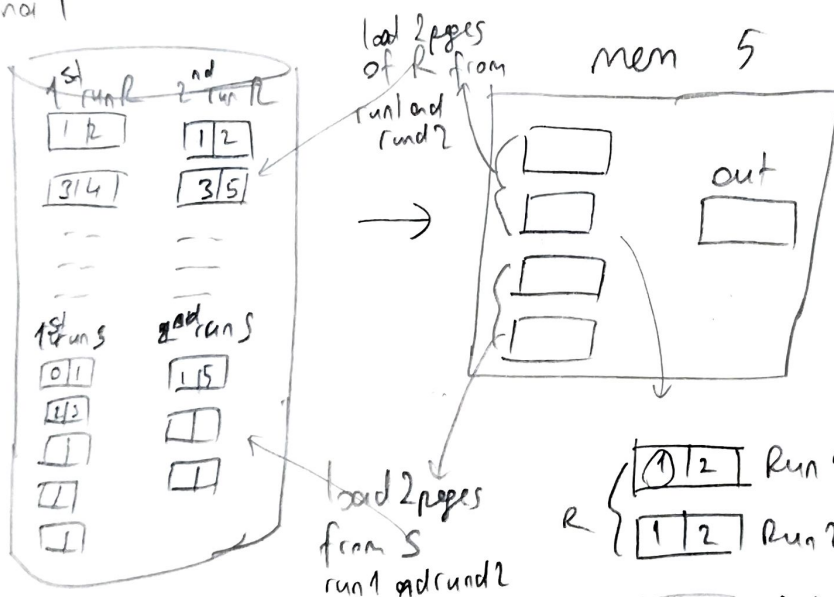
ex:
R has 10 pages, 2 tuples per page
S has 8 pages, //



do ~~this~~ this step for the remaining 5 pages of R.

do these steps for S too. $\Rightarrow 2B(R) + 2B(S)$

Final



$R \begin{cases} [1|2] \text{ Run 1} \\ [1|2] \text{ Run 2} \end{cases}$
 $S \begin{cases} [0|1] \text{ Run 1} \\ [1|5] \text{ Run 2} \end{cases}$
 output $\Rightarrow \begin{matrix} 1 & 1 \\ 1 & 1 \\ 1 & 1 \\ 1 & 1 \end{matrix}$

remove this page, it is finished

In our example total cost = $3B(R) + 3B(S)$

$R \begin{cases} [1|2] R1 \\ [1|2] R2 \end{cases} \rightarrow \begin{matrix} 2, 2 \\ 2, 2 \end{matrix}$
 $S \begin{cases} [2|3] R1 \\ [1|5] R2 \end{cases}$ continue

normal formula = $M1 = \frac{B(R)}{M}$ in our case $\frac{10}{5} = 2$


$M2 = \frac{B(S)}{M}$ in our case $\frac{8}{5} = 2$

if $M1 + M2 \leq M$, it can be done otherwise recursively (sort merge)

(7)

We can also do two pass based on hashing.

b) Hash



hash fonksiyonunun eşit bölünebilirliği diskinin, her bucket will have $\frac{B(R)}{M}$ size

if $\frac{B(R)}{M} \leq M$ i.e. $B(R) \leq M^2 \rightarrow$ each bucket will fit in memory.

Step 1: Hash into $M-1$ buckets
Send all buckets to disk


} Same hash function

Step 2: R --

Step 3: Join every pair of buckets } $h_2 \neq h_1$ (different hash)


Cost $3B(R) + 3B(S)$, assumption $\min(B(R), B(S)) \leq M^2$

example




Read S.

Step 1:



do this for all pages of S.


At the end, we will 4



They are sorted
but partitioned


Step 2: do the same step for R

Step 3: Using dif. hash funct.




$$\text{Cost} = 3B(R) + 3B(S)$$

① read 1st partition of R into M, h2 it.



② scan only matching partition of S.



Buraya R'leri hashleyerek
kaldığımız için, her gelen S'in
partitionu linear search yapmıyor! ✓

Daha hızlı!

from page

$0 \bmod 3 \rightarrow 0$	0, 12 not match	take next	<u>8, 12</u>
$4 \bmod 3 \rightarrow 1$	1, 4 //		$8 \bmod 3 \rightarrow 2$, 8, 8 ✓ match
	1, 4 //		$12 \bmod 3 \rightarrow 0$, 12, 12 ✓ match