

Ceng352 - Database Management Systems

Written Assignment 3

Spring 2019

Q1 Consider the following database schema:

```
Customer (cid, name, surname, phone, email)
Order (cid, pid, date, quantity)
Product (pid, pname, price)
```

For each of the queries below, show an equivalent relational algebra plan.

(a)

```
SELECT O.quantity, COUNT(DISTINCT C.cid),
FROM Customer C, Order O, Product P
WHERE C.cid = O.cid AND
      O.pid = P.pid AND
      C.email LIKE '%gmail.com' AND
      P.pid = 100
GROUP BY O.quantity
HAVING O.quantity > 10
```

(b)

```
SELECT C.name, C.phone
FROM Customer C
WHERE NOT EXISTS (
    SELECT *
    FROM Product P
    WHERE P.price < 50 AND NOT EXISTS (
        SELECT *
        FROM Order O
        WHERE O.cid = C.cid AND O.pid = P.pid
    )
)
```

Q2 Consider the following relations:

Professor (pid, name, dept) Teaching (pid, code, semester)

There are 1000 professors in 100 departments. The size of the relation **Professor** on disk is 200 pages. There are 10000 teaching records for 4 semesters. The size of the relation **Teaching** is 500 pages.

(a) Consider the following query:

<pre>SELECT * FROM Teaching T WHERE T.semester = 'F2017'</pre>
--

Estimate the cost of executing this query using

- Table scan
- A clustered B+ tree index on **semester**
- An unclustered B+ tree index on **semester**

(b) Consider the following query:

<pre>SELECT * FROM Professor P WHERE P.dept = 'CENG'</pre>
--

Estimate the cost of executing this query using

- Table scan
- A clustered hash index on **dept**
- An unclustered hash index on **dept**

(c) Consider the following query:

<pre>SELECT * FROM Professor P WHERE P.pid > 200 AND P.pid < 400</pre>
--

Estimate the cost of executing this query using

- Table scan
- A clustered B+ tree index on **pid** (assume that **pid** values range in between 0 and 1000)
- An unclustered B+ tree index on **pid** (assume that **pid** values range in between 0 and 1000)

(d) Consider the following query:

<pre>SELECT DISTINCT name FROM Professor</pre>
--

Assuming 10 names fit in one page, estimate the cost of executing this query using

- Hashing based projection (assume $M = 21$ available memory pages)
- Sorting based projection (assume $M = 20$ available memory pages)

Q3 Consider two relations $R(a, b)$ and $S(b, c)$ with the following statistics:

$T(R) = 10000$, $B(R) = 1000$ (each page contains 10 tuples),
 $V(R, b) = 200$ (number of distinct values of attribute b in R)

$T(S) = 6000$, $B(S) = 1500$ (each page contains 4 tuples),
 $V(S, b) = 150$ (number of distinct values of attribute b in S)

- (a) What is the cost of $R \bowtie S$ using **page-at-a-time nested loop join** algorithm?
- (b) What is the cost of $R \bowtie S$ using **block nested loop join** algorithm assuming that available memory $M = 102$ pages?
- (c) What is the cost of $S \bowtie R$ using **block nested loop join** algorithm assuming that available memory $M = 102$ pages?
- (d) Describe how to join the two relations if S has a **i)** clustered **ii)** unclustered index on the join attribute b . What is the cost of join using **index-nested loop join** for each type of index?
- (e) Describe in detail how to join these two relations using **partitioned hash join** given that 101 pages can fit in main memory at a time. Specially, explain how many pages are used in memory and how. Also specify what exactly is written to disk and when. Compute the cost of the partitioned hash join operation.
- (f) Describe in detail how to join these two relations using **merge-join** given that only 101 pages can fit in main memory at a time. Specially, explain how many pages are used in memory and how. Also specify what exactly is written to disk and when. Compute the cost of the merge-join operation.